

T-111.5360 Report:

Remote Mouse

Valter Kraemer 84669F
Ville Skyttä 42818N
Aalto University

December 28, 2015

1 Introduction

WebSockets is a technology providing an efficient full-duplex bidirectional communications channel between clients and servers. Due to its low latency characteristics, it is well suited for real time applications. The WebSocket protocol is standardized by IETF (Fette and Melnikov, 2011) and W3C (Hickson, 2012) develops an API to enable its use in web pages.

In this report we introduce Remote Mouse, a virtual mouse pointer on a web page, controlled from another web device. WebSockets is used as the communications technology between the devices, through an intermediate server.

[TODO: more]

2 Related work

Bassbouss et al. (2013) address multi-screen web application development and the transformation of traditional web applications to multi-screen capabilities. Both the current and proposed multi-screen application models utilize WebSocket in communications between clients (screens) and servers. [TODO: more about this and how it is related to our stuff]

Agar.io¹ is a massively multiplayer online game. In a nutshell, players control cells in a petri dish, attempting to grow larger by consuming pellets and other cells, and avoiding being consumed by other cells. The game is available for web browsers on its website, and Android and iOS versions are available for mobile devices. The web version uses an HTML canvas and its 2D context for rendering, as well as HTML animation frames. Communication between the browser and the server is implemented using WebSockets. Data is transferred using WebSocket binary frames (WebSocket opcode 2), which

are constructed using the ECMAScript 6 ArrayBuffer and DataView objects. Due to the binary nature of the data, the exact semantics of it are not available. During gameplay, the traffic consists of on the order of 50 WebSocket frames per second, with their sizes ranging approximately from a few to 200 bytes.

YouTube has a feature with which it is possible to control another YouTube window's video controls from another screen, such as a computer or a mobile device. It is primarily intended for controlling YouTube on smart TVs, but can also be used in a browser. The controlled screen² is operated by its paired remote³. The remote works by sending POST requests to the server that forwards them to the controlled window. It also uses polling every 10 seconds to check that the controlled device is still available. The controlled window uses long polling to check if any information is updated. YouTube is using LocalStorage for storing information about the playback device and different identifiers. MediaSource is used to attach sources to their video elements.

Remot.io⁴ is a service that controls HTML presentations such as reveal.js⁵ from touch based devices. The controlling device sends POST requests to the server that forwards them to the controlled device using long polling by the controlled device. Remot.io is using touch events for their remote. Swipe gestures translate to the directions the user wants to navigate in the slides.

3 Results

The controller is a web page. Most of it is empty, serving as the control area.

²<http://www.youtube.com/tv>

³<http://www.youtube.com/pair>

⁴<http://remot.io/>

⁵<http://lab.hakim.se/reveal-js>

¹<http://agar.io/>

For non-touch devices, the document has mouse-move and click event handlers. Each mousemove event results in the pointer position being sent to the server in a WebSocket message.

For touch devices, the document contains event listeners for touchstart, touchmove, and touchend.

Positions are sent as pairs of floating point numbers representing the relative left and top position offsets of the event in percents, proportional to the dimensions of the controller window. The controllee moves the virtual pointer according to these values, by using them as percentages for the virtual pointer’s left and top CSS properties.

WebSocket messages for click events contain only the information that the controller requests a click; there is no position information included with it. The controllee locates the element which is the target of the click based on its own virtual pointer position, using the `Document.elementFromPoint` CS-SOM View Module method⁶.

4 Analysis

[TODO]

The most important technology for both controller and controllee sides of the Remote Mouse implementation is WebSockets. According to caniuse.com, it is fully supported in all current major browsers since 2013⁷. The Node.js ws library used on the server side has had releases available from GitHub since 2011⁸.

The DeviceOrientation API used to implement scrolling based on controller orientation is also well supported in current browsers to the extent required by Remote Mouse. According to caniuse.com⁹, only Microsoft Edge has full support for it, most other browsers have partial support, and the desktop version of Safari has none. Safari’s non-support is not a major problem, because controller devices are expected to be mobile ones, and the iOS Safari supports the API.

Table 1 summarizes support for WebSockets and DeviceOrientation in major browsers. The versions listed are the first ones in which full support for the technology appeared, followed by the release year in parenthesis. If full support is not yet available, the version number in parenthesis indicates the first version with partial support.

To test the subjective effect of latency on user experience, a test with three users was conducted. The

	WebSockets	DeviceOrientation
IE	10 (2012)	(11) (2013)
Edge	12 (2015)	(12) (2015)
Firefox	11 (2012)	(6) (2011)
Chrome	16 (2011)	(7) (2010)
Safari	7 (2013)	-
iOS Safari	6.1 (2013)	(4.3) (2011)
Android Browser	4.4 (2013)	(3) (2011)
Chrome for Android	47 (2015)	(47) (2015)

Table 1: caniuse.com: WebSockets and DeviceOrientation support in selected browsers

Latency	Grade	Latency	Grade	Latency	Grade
User 1		User 2		User 3	
500 ms	2	400 ms	2	0 ms	8
200 ms	4	100 ms	5	500 ms	1
100 ms	5	300 ms	4	300 ms	4
0 ms	6	0 ms	8	100 ms	8
100 ms	6	100 ms	8	0 ms	9
0 ms	8	0 ms	8	200 ms	7

Table 2: Subjective user test results

users were first asked to use an Apple Magic Trackpad¹⁰ to get a feeling of a local, low latency user experience. Then, they were asked to use Remote Mouse with the latency throttle set to varying settings. The latency settings were shuffled, i.e. not presented in increasing or decreasing order in order to avoid users’ expectations affecting the results. Users were tasked to grade the quality of the pointer control experience in scale from 0 to 10, with grade 0 being the lowest one, equal to unusable, and 10 being equally good as the Magic Trackpad.

All three users rated the experience to belong in the middle of the scale at approximately 250 ms latency. 500 ms was classified as barely usable, and 0 to 100 ms quite acceptable. User test data is included in table 2.

To aid in estimating how these estimates translate to use of Remote Mouse in different network setups, table 3 lists the typical latencies when the service is running in Heroku and locally, and when it is being used over different network connections.

[TODO: something about bandwidth usage, single WS message per MTU?]

⁶<http://www.w3.org/TR/cssom-view/>

⁷<http://caniuse.com/#feat=websockets>

⁸<https://github.com/websockets/ws/releases>

⁹<http://caniuse.com/#feat=deviceorientation>

¹⁰https://en.wikipedia.org/wiki/Magic_Trackpad

Network	Server	Latency
2G	Heroku	500 ms
3G	Heroku	70 ms
LTE	Heroku	70 ms
Wi-Fi	Heroku	70 ms
Wi-Fi	local	5 ms

Table 3: Typical setup latencies

5 Conclusions

[TODO]

Latency of a network connection is much more important for satisfactory user experience with Remote Mouse than its bandwidth. Bandwidth needs of the application are already quite modest with the current implementation, and could be further reduced, for example by using a more efficient binary WebSocket message payloads, and compression. However, given the already low requirements and possibility of getting negative effects on latency from optimizing for bandwidth usage, these possibilities were not pursued as they are not likely to result in significant overall user experience improvements, if any.

Based on the test conducted as well as the authors' own experiences, the latency goal for acceptable Remote Mouse user experience should be set to the 0 to 100 ms range. According to our test results, these kinds of latencies can be achieved with 3G and better mobile network connections; 2G connectivity is not sufficient.

The technology stack related to WebSockets is stable and ready for production use in both browser and server side. We did not run into any issues on either browser or server side during the Remote Mouse development process that would have been related to WebSockets implementations. On the contrary, we found the APIs and implementations very easy to use, and their performance matches or exceeds the requirements for Remote Mouse.

A prominent use case for Remote Mouse is remote control of web based presentations and applications, using for example a laptop computer to host the presentation or web application and displaying its screen to viewers, while controlling it remotely from a mobile device. Because the contents of the controllee screen are not visible in the controller, this use case is in our implementation limited to setups where the user operating the controller can see the controllee screen. If the controllee screen would be available for remote viewing, use cases like for example remote assistance of web application use would be quite relevant. The technology and principle of tracking the pointer or touch movements could also be used for recording user actions on a

web site, for example for usability evaluation, user interface research, and trials.

References

- Louay Bassbouss, Marc Tritschler, Stephan Steglich, Kiyoshi Tanaka and Yuji Miyazaki. Towards a Multi-screen Application Model for the Web. *Computer Software and Applications Conference Workshops (COMPSACW), 2013 IEEE 37th Annual*, pages 528–533. IEEE, July 2013. doi: 10.1109/COMPSACW.2013.96.
- Ian Fette and Alexey Melnikov. The WebSocket Protocol. RFC 6455, December 2011. URL <http://www.ietf.org/rfc/rfc6455.txt>.
- Ian Hickson. The WebSocket API. Technical Report, W3C, September 2012. URL <http://www.w3.org/TR/2012/CR-websockets-20120920/>.